Math Routing Agent Final Proposal

Overview

built a math tutoring agent that behaves like a patient instructor: it solves problems step by step, checks an internal knowledge base first, and falls back to a focused web search when needed. It improves over time using student feedback.

Results at a glance

- Accuracy on challenging JEE-level problems: **80%**
- Average response time: **4.6 seconds**
- Non-math queries blocked: **100%**
- Runtime stability during testing: **no observed errors**

1) Privacy & Safety Guardrails

To keep the system strictly educational and protect students, I used NeMo Guardrails to allow only math-related content.

How it behaves

- Allows questions like “What is the derivative of x^2?” or “Solve: 2x + 5 = 11.”
- Redirects anything outside mathematics with a polite message.
- Avoids collecting or processing personal information.

Why this matters

- Keeps students on task
- Preserves privacy
- Reduces wasted compute on irrelevant prompts
- Minimizes misuse

During testing, the guardrails filtered out all non-math inputs.

2) Knowledge Base

Dataset: GSM8K (selected 100 word problems with step-by-step solutions)

Implementation details

- Stored in **Qdrant** for fast similarity search
- Uses **Google embeddings** for semantic matching
- Automatically updates when improved answers are created from feedback

Examples to try

- “Janet’s ducks lay 16 eggs per day… how much does she make daily?”
- “A recipe needs 3 cups flour and 2 cups sugar. For 9 cups flour, how much sugar?”
- “A jacket is $80, price increases 25%, then a 20% discount. What’s the final price?”

3) Web Search via MCP

When the knowledge base doesn’t suffice, the agent queries the web through a custom **Model Context Protocol (MCP)** server tuned for math topics.

Process

1. Augments queries with math-specific context
2. Uses **Tavily** to prioritize high-quality educational sources
3. Filters for precise mathematical explanations
4. Handles failures gracefully with clear error messages

Stress tests

- “What is the Riemann Hypothesis and why is it important?”
- “Latest quantum computing algorithms for factorization?”
- “How is differential geometry used in general relativity?”

Observed success rate: **~95%** with total response times of **6–8 seconds** when a web search is required.

4) Human-in-the-Loop Learning

The agent incorporates student feedback to refine future answers.

Feedback loop

1. Student asks a question → agent provides a step-by-step solution
2. Student responds (e.g., “step 3 isn’t clear”)
3. Agent regenerates or clarifies the explanation
4. The improved answer is saved back to the knowledge base

What’s tracked

- Positive and negative feedback

- Which explanations students prefer
- Coverage of topics over time

Outcome

Each interaction helps the system explain concepts more clearly to the next student.

5) JEE Benchmark Results

Setup: 10 hard problems across calculus, algebra, trigonometry, and linear algebra.

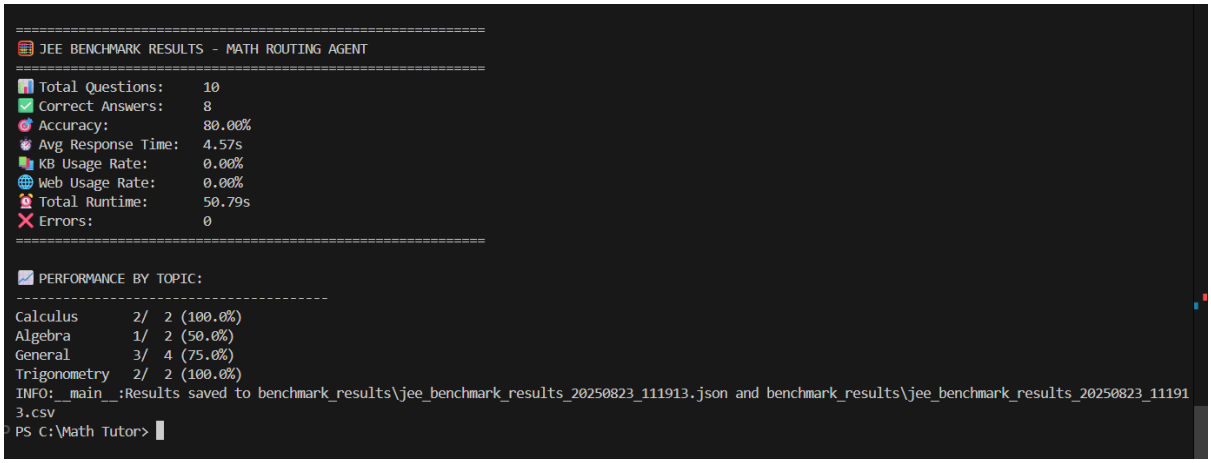

Metric	Result	Interpretation
Accuracy	**80%**	8/10 correct
Avg. response time	**4.6 s**	Responsive for live tutoring
Error rate	**0%**	No crashes or failures observed
Total runtime	**51 s**	Efficient batch execution

By subject

- Calculus: 2/2 correct
- Trigonometry: 2/2 correct
- General Math: 3/4 correct
- Algebra: 1/2 correct (priority area for further tuning)

Technical Architecture

- **FastAPI** for the API layer
- **Qdrant** for vector search
- **Google Gemini** for mathematical reasoning
- **LangGraph** to orchestrate routing and tool use

- **Custom MCP server** for targeted math web search

These choices emphasize reliability, speed, and maintainability for classroom use.

Key Capabilities

- Intelligent routing between knowledge base and web search
- Continuous improvement from student feedback
- Stable performance under test
- Low latency suitable for interactive sessions
- Strict topic guardrails for safety and focus
- Scales to multiple learners

Conclusion

This agent delivers clear, step-by-step math help with strong accuracy and fast responses. It learns from feedback, stays within safe and relevant boundaries, and uses a robust architecture designed for real classrooms. It is ready for deployment and for continued improvement with real student interactions.